

Unit 5


LASSO Regularization and Feature Selection

EE-UY 4563/EL-GY 9143: INTRODUCTION TO MACHINE LEARNING
PROF. SUNDEEP RANGAN (WITH MODIFICATION BY YAO WANG)

Learning Objectives

- ☐ Describe **model selection** and identify when it may be needed
- ☐ Mathematically describe linear regression with **regularization**
- ☐ Select regularizers to impose constraints such as sparsity
- ☐ Compute an L1-regularized estimate (LASSO) using sklearn tools
- ☐ Compute the optimal regularization level using cross validation
- ☐ Interpret results from a LASSO path
- ☐ Set regularizer based on a probabilistic prior
- ☐ Perform other feature selection methods

Outline

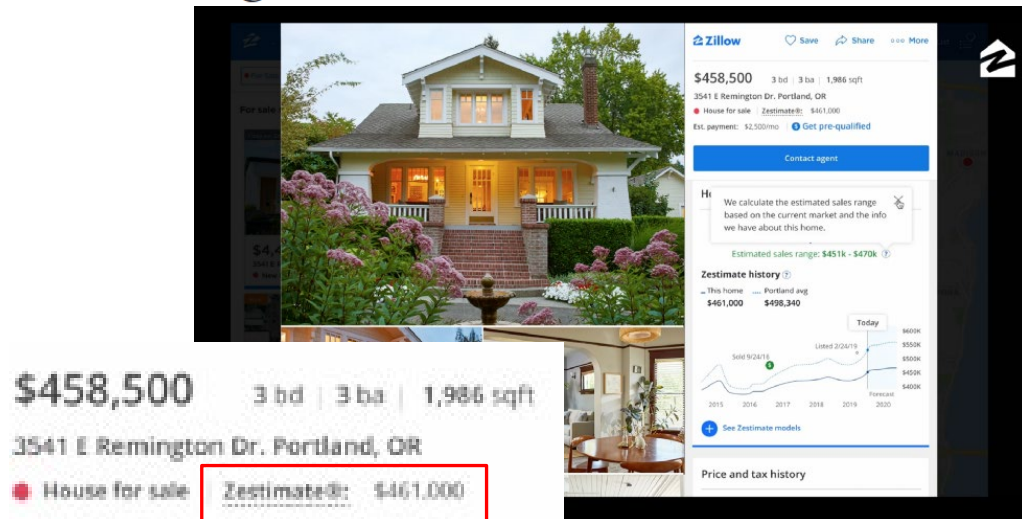
 Motivating Example: Feature selection in predicting housing prices

- ☐ Model selection and regularization
- ☐ Housing prices prediction with LASSO
- ☐ Probabilistic interpretation
- ☐ Other model selection examples
- ☐ Other model selection methods
- ☐ In-Class Exercise: Audio Pitch Detection

Predicting Housing Prices

AI, MACHINE LEARNING & RESEARCH

Introducing a new and improved Zestimate algorithm



<https://www.zillow.com/tech/introducing-a-new-and-improved-zestimate-algorithm/>

❑ Many services now predict house prices

❑ Data science enters real estate!

❑ Many possible variables:

- Square meters
- Condition
- Zip code
- Education quality
- ...

❑ What variables *really* determine the price?



Ames, Iowa Dataset



Ames, Iowa: Alternative to the Boston Housing Data as an End of Semester Regression Project

[Dean De Cock](#)

Truman State University

Journal of Statistics Education Volume 19, Number 3(2011),
www.amstat.org/publications/jse/v19n3/decock.pdf

Copyright © 2011 by Dean De Cock all rights reserved. This text may be freely shared among individuals, but it may not be republished in any medium without express written consent from the author and advance notification of the editor.

Key Words: Multiple Regression; Linear Models; Assessed Value; Group Project.

Abstract

This paper presents a data set describing the sale of individual residential property in Ames, Iowa from 2006 to 2010. The data set contains 2930 observations and a large number of explanatory variables (23 nominal, 23 ordinal, 14 discrete, and 20 continuous) involved in assessing home values. I will discuss my previous use of the Boston Housing Data Set and I will suggest methods for incorporating this new data set as a final project in an undergraduate regression course.

- ❑ Ames, Iowa Dataset
 - Sales from 2006 to 2010
 - From Dean De Cock
- ❑ Alternative to Boston Housing dataset
- ❑ Many more variables to explore
 - Approximately 81 variables
 - 2930 samples



NYU

**TANDON SCHOOL
OF ENGINEERING**



Loading the Dataset

```
1 df = pd.read_csv('housing_train.csv')
2 df.head()
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0

5 rows × 81 columns

Issues:

- Many different types of data: Discrete and continuous
- Missing values (NaN)

Data Cleaning

```
1 nsamp, natt = df.shape
2 print('Number samples = %d' % nsamp)
3 print('Number attributes per sample = %d' % natt)
```

Number samples = 1460
Number attributes per sample = 81

```
1 df = df.dropna(axis=1)
2
3 nsamp, natt = df.shape
4 print('Number samples = %d' % nsamp)
5 print('Number attributes per sample = %d' % natt)
6
```

```
1 df = df.loc[df['SaleCondition'] == 'Normal']
2
3 nsamp, natt = df.shape
4 print('Number samples = %d' % nsamp)
5 print('Number attributes per sample = %d' % natt)
```

Number samples = 1198
Number attributes per sample = 62

❑ Original data

❑ Remove columns with NaN values

- Could use more sophisticated methods

❑ Keep only normal sales

- Recommended in De Cock paper
- Makes fitting much easier

Categorical Variables

- ❑ Data has many categorical variables
- ❑ Need to code the categorical variables to numerical values

Real valued

Categorical

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	0
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	0
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0

Splitting the Variables

❑ First, split the variables into categorical and real

```
# Remove the ID, month sold and sales price (it is the target)
ignore_vars = ['Id', 'MoSold', 'SalePrice']

# Find real and categorical variables
cols = df.columns
cat_vars = []
real_vars = []

for col in cols:
    if not (col in ignore_vars):
        if df.dtypes[col] == 'object':
            cat_vars.append(col)
        else:
            real_vars.append(col)
```

Categorical variables = ['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2', 'Heating', 'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual', 'Functional', 'FireplaceQu', 'GarageType', 'GarageFinish', 'GarageQual', 'GarageCond', 'PavedDrive', 'PoolQC', 'Fence', 'MiscFeature', 'SaleType', 'SaleCondition']

Real variables = ['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces', 'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal', 'YrSold']

One Hot Coding

Original

	MSZoning	Street	LotShape	LandContour	Utilities	LotConfig
0	RL	Pave	Reg	Lvl	AllPub	Inside
1	RL	Pave	Reg	Lvl	AllPub	FR2
2	RL	Pave	IR1	Lvl	AllPub	Inside
4	RL	Pave	IR1	Lvl	AllPub	FR2
5	RL	Pave	IR1	Lvl	AllPub	Inside

One Hot coded

	MSZoning_FV	MSZoning_RH	MSZoning_RL	MSZoning_RM	Street_Pave
0	0	0	1	0	1
1	0	0	1	0	1
2	0	0	1	0	1
4	0	0	1	0	1
5	0	0	1	0	1

- ❑ Use pandas `get_dummies`
- ❑ Replaces categorical variables with one-hot coded values
 - Ex: MSZoning
 - Becomes MSZoning_FV, MSZoning_RH, ...

```
# Get the dataframes with real and categorical variables
df_real = df[real_vars]
df_cat = df[cat_vars]

# One-hot encode the categorical variables
df_cat_enc = pd.get_dummies(df_cat, drop_first=True)
```

Scaling Data

- ❑ Split data into training and test
- ❑ Scale data
 - Remove mean and variance
- ❑ Needed to compare coefficients
 - Ensures that all variables have same range
- ❑ Note: The scaling transform is
 - Fit on the training data
 - Performed on training and test

```
from sklearn.model_selection import train_test_split  
  
Xtr, Xts, ytr, yts = train_test_split(X,y,test_size=0.3)
```

```
from sklearn.preprocessing import StandardScaler  
  
# Create the scaler objects  
xscal = StandardScaler()  
yscal = StandardScaler()  
  
# Fit and transform the training data  
Xtr1 = xscal.fit_transform(Xtr)  
ytr1 = yscal.fit_transform(ytr[:,None])  
  
# Transform the test data  
Xts1 = xscal.transform(Xts)  
yts1 = yscal.transform(yts[:,None])
```

First Try: Linear Regression

```
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.metrics import r2_score

# Fit
reg_ls = LinearRegression()
reg_ls.fit(Xtr1, ytr1)

# Training R^2
yhat1_tr = reg_ls.predict(Xtr1)
rsq_tr = r2_score(yhat1_tr, ytr1)
print('Training R^2 = %12.4e' % rsq_tr)

# Test R^2
yhat1_ts = reg_ls.predict(Xts1)
rsq_ts = r2_score(yts1, yhat1_ts)
print('Test R^2      = %12.4e' % rsq_ts)
```

```
Training R^2 = 9.3726e-01
Test R^2      = -1.0430e+20
```

- ❑ Simple idea:
 - Use linear regression over features
- ❑ Fits the training data very well!
 - $R^2 \approx 0.937$
- ❑ But, completely fails on the test data
 - $R^2 > 10^{20}$

Conditioning

- ❑ What went wrong?
- ❑ Recall LS solution is: $\hat{\beta} = (A^T A)^{-1} A^T y$
- ❑ Matrix $A^T A$ may be ill-conditioned
 - Eigenvalues close to zero
 - Inverse blows up
- ❑ With ill-conditioned data:
 - Training error is fine
 - But the test error blows up
- ❑ Overfits data

```
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.metrics import r2_score

# Fit
reg_ls = LinearRegression()
reg_ls.fit(Xtr1, ytr1)

# Training R^2
yhat1_tr = reg_ls.predict(Xtr1)
rsq_tr = r2_score(yhat1_tr, ytr1)
print('Training R^2 = %12.4e' % rsq_tr)

# Test R^2
yhat1_ts = reg_ls.predict(Xts1)
rsq_ts = r2_score(yts1, yhat1_ts)
print('Test R^2      = %12.4e' % rsq_ts)
```

```
Training R^2 = 9.3726e-01
Test R^2      = -1.0430e+20
```

Improving Conditioning via Ridge Regression

❑ Standard LS solution: $\hat{\beta} = (A^T A)^{-1} A^T y$

❑ **Ridge Regression**: Add a conditioning term:

$$\hat{\beta} = (A^T A + cI)^{-1} A^T y$$

- c is a small positive value.
- Makes inverse well-behaved
- We will see this technique more later

❑ Get good test R^2

```
reg_ls = Ridge(alpha=1e-5)
reg_ls.fit(Xtr1, ytr1)
yhat1 = reg_ls.predict(Xts1)
rsq = r2_score(yts1, yhat1)
print('Test R^2      = %12.4e' % rsq)
```

Test R^2 = 0.904567

What Components Matter?

❑ Simple idea: Look at large coefficients

❑ We see variables that we may expect:

- Square footage
- Quality
- Zoning

❑ But there are some issues

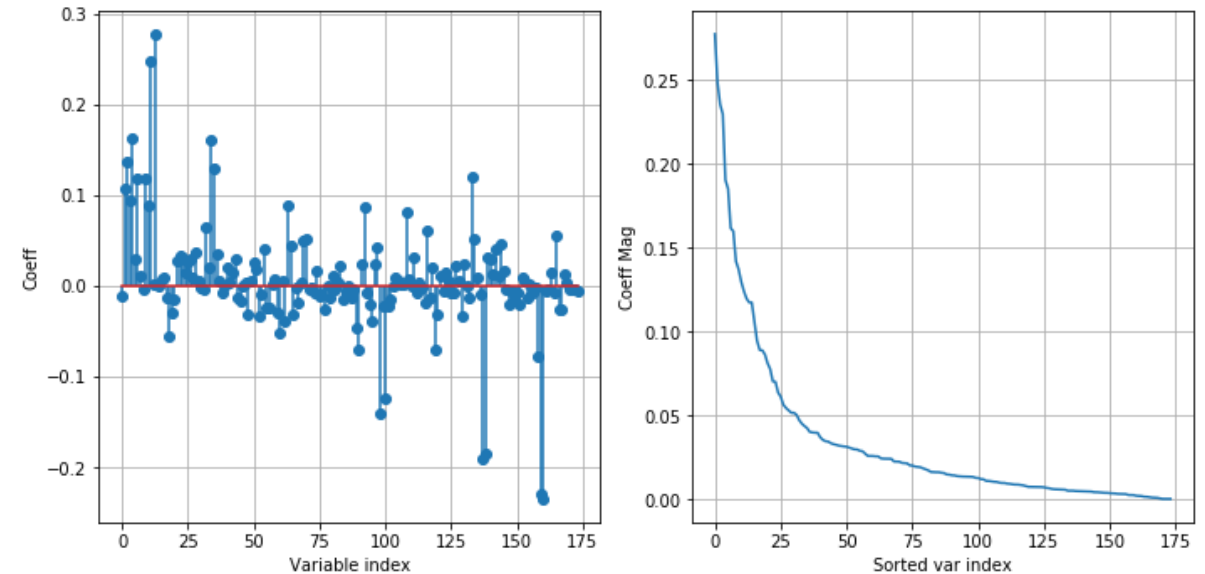
- Some variables seem highly correlated
- Ex: GrLivArea and 2ndFlrSF
- KitchenQual and OverallQual
- Do we need both?

```
1 coeff_ls = reg_ls.coef_.ravel()
2 nprint = 10
3 I = np.argsort(np.abs(coeff_ls))
4 I = np.flipud(I)
5 for i in range(nprint):
6     j = I[i]
7     print('%20s %f' % (xnames[j], coeff_ls[j]))
```


```
GrLivArea 0.277146
2ndFlrSF 0.248113
KitchenQual_TA -0.235208
KitchenQual_Gd -0.229472
ExterQual_Gd -0.190125
ExterQual_TA -0.184943
YearBuilt 0.161462
MSZoning_RL 0.159591
RoofStyle_Gable -0.141669
OverallQual 0.136913
```

What Components Do *Not* Matter?

- ❑ All coefficients are far from zero
- ❑ Very few coefficients that can be removed
- ❑ Does this mean all variables matter?
- ❑ Model or feature selection problem:
 - *How do we find the variables that matter?*



Outline

- ☐ Motivating Example: Feature selection in predicting housing prices
-  ☐ Model selection and regularization
 - ☐ Housing prices prediction with LASSO
 - ☐ Probabilistic interpretation
 - ☐ Other model selection examples
 - ☐ Other model selection methods
- ☐ In-Class Exercise: Audio Pitch Detection

Model Selection via Sparsity

	MSZoning_FV	MSZoning_RH	MSZoning_RL	MSZoning_RM	Street_Pave	LotShape_IR2	LotShape_IR3	LotShape_Reg	LandContour_HLS
0	0	0	1	0	1	0	0	1	0
1	0	0	1	0	1	0	0	1	0
2	0	0	1	0	1	0	0	0	0
4	0	0	1	0	1	0	0	0	0
5	0	0	1	0	1	0	0	0	0

174 variables after
one-hot coding

❑ **Model selection problem:** Need to identify the parameters that *really* matter

- Help interpret results
- Improves generalization error (less parameters)

❑ **Idea:** Fit model under **sparsity constraint**:

- Linear model: $\hat{y} = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$
- Feature x_j is ignored if $\beta_j = 0$
- Try to force most $\beta_j = 0 \Rightarrow$ Model only uses a few of the variables

Regularized LS Estimation

❑ **Regularization:** General method for finding constrained solutions

- E.g. solutions that are sparse

❑ Standard least squares estimation (from Lecture 3):

$$\hat{\beta} = \arg \min_{\beta} MSE(\beta), \quad MSE(\beta) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

❑ **Regularized estimator:**

$$\hat{\beta} = \arg \min_{\beta} J(\beta), \quad J(\beta) = MSE(\beta) + \phi(\beta)$$

- $MSE(\beta)$ = mean-squared prediction error from before
- $\phi(\beta)$ = regularizing function.

❑ **Concept:** Regularizer penalizes β that are “unlikely”

- Constrains estimate to smaller set of parameters

Two Common Regularizers

❑ Ridge regression (called L2)

$$\phi(\beta) = \frac{\alpha}{n} \sum_{j=1}^d |\beta_j|^2$$

❑ LASSO regression (called L1)

$$\phi(\beta) = 2\alpha \sum_{j=1}^d |\beta_j|$$

❑ Coefficient $\alpha > 0$ determines regularization level

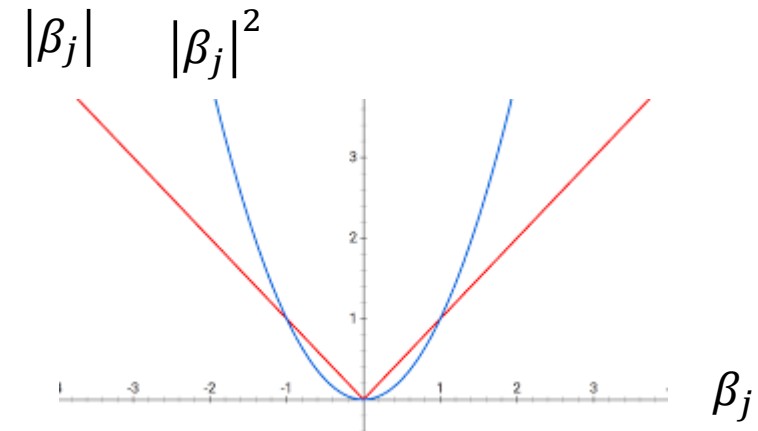
- Higher $\alpha \Rightarrow$ Higher level of regularization, more constrained
- Will show how to select α later via cross-validation
- Scaling factors adjust to match sklearn convention

❑ Both penalize large β_j : Tries to make β_j small

- Will see that L1 also promotes sparsity

❑ Convention: Do not include intercept term β_0

- In general, no reason to make this term small



L1 and L2 Norm

□ Ridge and LASSO Regularization can be written with **norms**

□ **Ridge** cost function:

$$J(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha \sum_{j=1}^d |\beta_j|^2 = \|\mathbf{y} - \mathbf{A}\boldsymbol{\beta}\|^2 + \alpha \|\boldsymbol{\beta}\|_2^2$$

□ **LASSO** cost function:

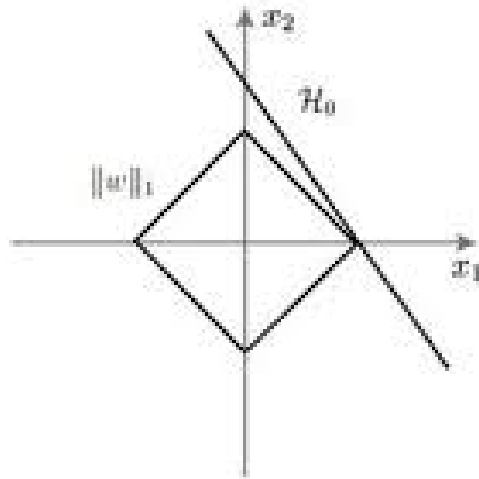
$$J(\boldsymbol{\beta}) = \frac{1}{2n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha \sum_{j=1}^d |\beta_j| = \frac{1}{2n} \|\mathbf{y} - \mathbf{A}\boldsymbol{\beta}\|^2 + \alpha \|\boldsymbol{\beta}\|_1$$

◦ $\|\boldsymbol{\beta}\|_1$ = **L1 norm** (pronounced ell-1)

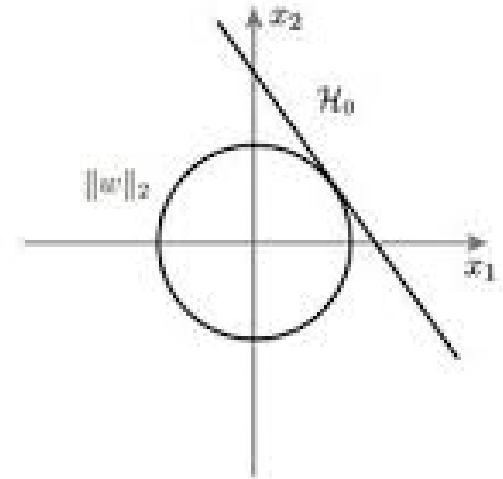
Ridge vs LASSO

- ❑ L2 tends to lead to many “small” coefficients
 - But solutions are not exactly zero
 - Not ideal for feature selection
- ❑ L1 tends to lead to more **sparse** solutions
 - Several coefficients are zero

A L1 regularization



B L2 regularization



Solving Ridge Regression

- Ridge regression problem: Find β to minimize

$$J(\beta) = \|\mathbf{y} - A\beta\|^2 + \alpha\|\beta\|^2$$

- Solution for given regularization level

$$\beta_{ridge} = (A^T A + \alpha I)^{-1} A^T \mathbf{y}$$

- Set gradient = 0
- See homework

- Sklearn function for ridge regression:

- http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Ridge.html

Solving LASSO Regression

□ LASSO cost function:

$$J(\boldsymbol{\beta}) = \frac{1}{2n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha \sum_{j=1}^d |\beta_j| = \frac{1}{2n} \|\mathbf{y} - \mathbf{A}\boldsymbol{\beta}\|^2 + \alpha \|\boldsymbol{\beta}\|_1$$

□ Because derivative of $|\beta_j|$ is not continuous, there is no closed-form solution.

□ Many methods to solve iteratively

- Least angle regression (LAR), coordinate descent, ADMM
- However, the cost function is convex \Rightarrow no local minima [See Unit 7]
- Beyond the scope of this class
- See textbook [Hastie2008] for LAR method

Data Scaling

❑ Scaling: Whenever using regularization:

- Scale each feature and the target to have zero mean and unit variance (or STD)
- $x_{ij} \rightarrow (x_{i,j} - \bar{x}_j) / \text{STD}(x_{ij})$
- $y_i \rightarrow (y_i - \bar{y}) / \text{STD}(y_i)$

❑ After predictor for the scaled data are determined:

- Derive the equivalent predictor on the original data (HW!)

❑ Motivation:

- Without scaling, the regularization level depends on the data range
- With mean removal, we do not need the intercept term β_0
- So that the regularization term is simply a L2 or L1 norm of coefficient vector

Selecting the Regularization Level

- How do we select regularization level α ?
 - Higher $\alpha \Rightarrow$ More constrained / simpler model
 - Lower $\alpha \Rightarrow$ More complex model
- Similar to inverse of model order
- Find α via cross-validation

Pseudo-code

Split in training X_{tr}, y_{tr} and test X_{ts}, y_{ts} .

For α in α_{test} :

- $\hat{\beta} = \text{fit}(X_{tr}, y_{tr}, \alpha)$ // Fit on training data
- $\hat{y}_{ts} = \text{predict}(X_{ts})$ // Predict on test data
- $S[\alpha] = \text{score}(y_{ts}, \hat{y}_{ts})$ // Score on test data

$\hat{\alpha} = \text{argmax } S[\alpha]$ // Select α with highest test score

Summary

Method	Regularizer	Effect on parameters	Solution for Fitting
None	$\phi(\boldsymbol{\beta}) = 0$	Leaves parameters unconstrained	$\hat{\boldsymbol{\beta}} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$
Ridge	$\phi(\boldsymbol{\beta}) = \frac{\alpha}{n} \ \boldsymbol{\beta}\ _2^2$	Makes parameters small Close to zero	$\hat{\boldsymbol{\beta}} = (\mathbf{A}^T \mathbf{A} + \alpha \mathbf{I})^{-1} \mathbf{A}^T \mathbf{y}$
LASSO	$\phi(\boldsymbol{\beta}) = 2\alpha \ \boldsymbol{\beta}\ _1$	Makes parameters sparse. Many coefficients exactly zero	No analytic solution. Need to run an optimizer

□ Regularized least squares

$$\hat{\boldsymbol{\beta}} = \arg \min_{\boldsymbol{\beta}} J(\boldsymbol{\beta}), \quad J(\boldsymbol{\beta}) = \frac{1}{n} \|\mathbf{y} - \mathbf{A}\boldsymbol{\beta}\|^2 + \phi(\boldsymbol{\beta})$$

□ Whatever you choose for the regularizer:

- Scale data before training
- Select regularization level with cross-validation

In-Class Exercise

Question

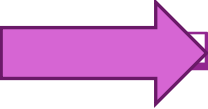
We wish to fit a linear model of the form,

$$\hat{y} = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p,$$

for a problem with $p = 100$ variables. Suggest regularizers $\phi(\boldsymbol{\beta})$ to impose the following constraints:

- (a) All the coefficients $\beta_j, j = 1, \dots, p$ should be close to zero, but not necessarily exactly zero.
- (b) Most of the coefficients $\beta_j, j = 1, \dots, p$ should be exactly zero
- (c) Among the first fifty coefficients, $\beta_j, j = 1, \dots, 50$, most coefficients should be zero. But, the other coefficients should be unconstrained.

Outline

- ❑ Motivating Example: Feature selection in predicting housing prices
- ❑ Model selection and regularization
- ❑ Housing prices prediction with LASSO
- ❑ Probabilistic interpretation
- ❑ Other model selection examples
- ❑ Other model selection methods
- ❑ In-Class Exercise: Audio Pitch Detection

LASSO Regression in Python

❑ Sklearn built in Lasso class

❑ Easy to use

- Set alpha
- Fit on training data
- Predict and score on test

Test $R^2 = 0.899122$

```
1 from sklearn.linear_model import Lasso
2 from warnings import simplefilter
3 from sklearn.exceptions import ConvergenceWarning
4 simplefilter("ignore", category=ConvergenceWarning)
5
6 # Select alpha
7 alpha = 3e-3
8
9 # Create Lasso object and fit on training data
10 reg = Lasso(alpha=alpha)
11 reg.fit(Xtr1, ytr1)
12
13 # Predict and score on test
14 yhat1 = reg.predict(Xts1)
15 rsq = r2_score(yts1, yhat1)
16
17 print('Test  $R^2 =$  %f' % rsq)
```

Optimizing Alpha via Cross Validation

- ❑ In each fold we:
 - Split data into training and test
 - Fit the scale on the training
 - Transform training and test
 - For each alpha:
 - Fit training and score on test
- ❑ Note: Scaling is redone on each fold
 - Ensures scaling is part of the training

```
10 # Run the cross-validation
11 rsq = np.zeros((nalpha, nfold))
12 for ifold, ind in enumerate(kf.split(X)):
13
14     # Get the training data in the split
15     Itr, Its = ind
16     Xtr = X[Itr, :]
17     ytr = y[Itr]
18     Xts = X[Its, :]
19     yts = y[Its]
20
21     # Fit and transform the data
22     Xtr1 = xscal.fit_transform(Xtr)
23     Xts1 = xscal.transform(Xts)
24     ytr1 = yscal.fit_transform(ytr[:, None])
25     yts1 = yscal.transform(yts[:, None])
26
27     for i, alpha in enumerate(alphas):
28
29         # Fit on the training data
30         reg = Lasso(alpha=alpha)
31         reg.fit(Xtr1, ytr1)
32
33         # Score on the test data
34         yhat1 = reg.predict(Xts1)
35         rsq[i, ifold] = r2_score(yts1, yhat1)
36
37     print('Fold = %d' % ifold)
38
39 # Compute mean and SE
40 rsq_lasso_mean = np.mean(rsq, axis=1)
41 rsq_lasso_se = np.std(rsq, axis=1) / np.sqrt(nfold-1)
```

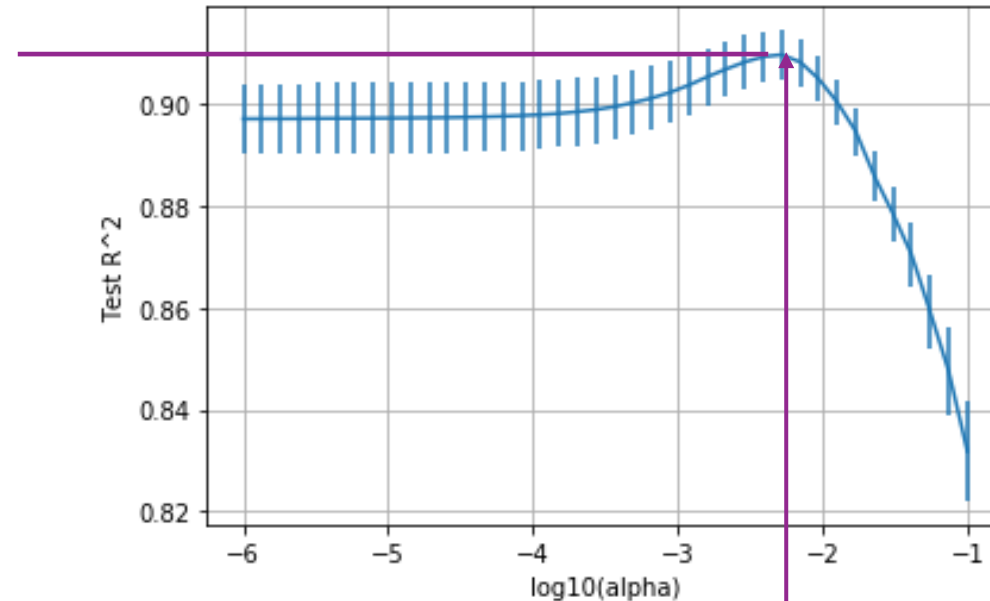
Cross Validation: Normal Rule

Max mean test R^2

□ Select alpha to maximize mean test R^2

- Normal rule
- Lower values of $\alpha \Rightarrow$ overfit
- Higher values of $\alpha \Rightarrow$ underfit

Alpha optimal (normal rule) = $5.2233e-03$
Mean test R^2 (normal rule) = 0.910



Under-regularized
Overfit

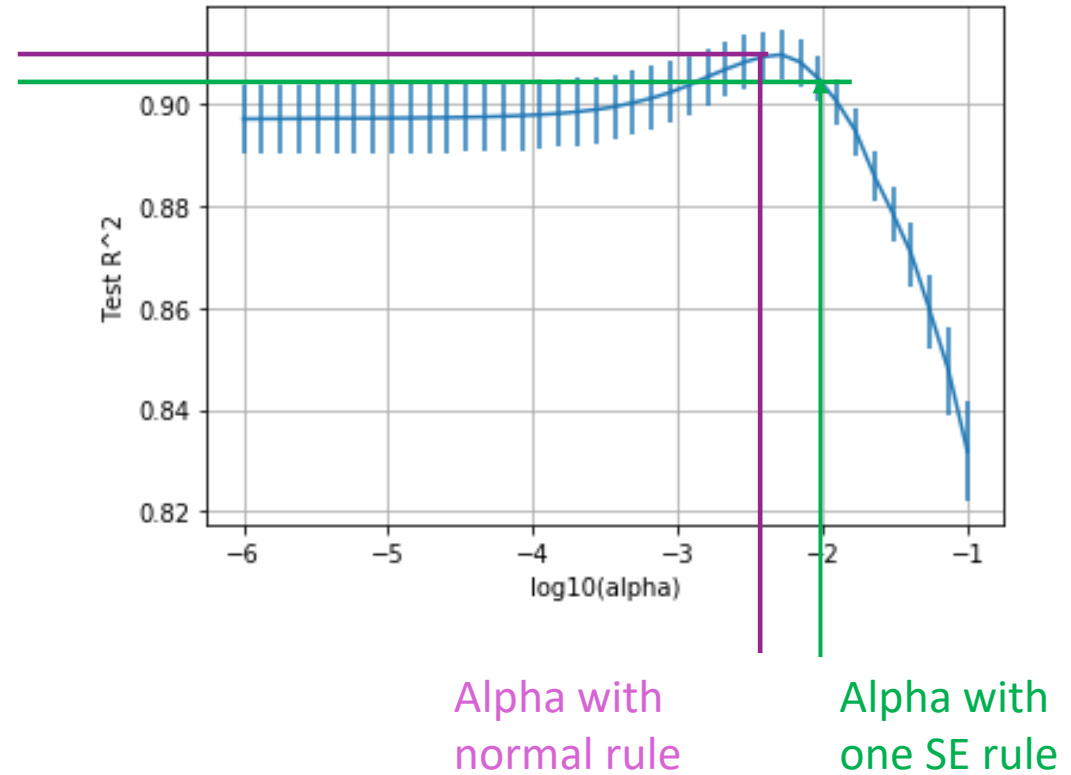
Optimal Alpha
(normal rule)

Over-regularized
Underfit

Cross Validation: One SE Rule



Max mean test R^2
Max mean test R^2 -one SE

- Can also use one SE rule:
 - Selects a higher regularized model
 - More sparse solution



Sparsity in the Coefficients

- Adding L1 regularization:
 - Makes coefficient smaller
 - Many coefficients approx. 0



Most Important Variables

❑ Right table: Variables with 10 large coefficient magnitudes

❑ Minimally regularized (Ridge) has:

- Variables that are highly correlated
- Ex: GrLivArea and 2ndFlrSF
- Several large variables


❑ Lasso:

- Reduces correlated variables
- Selects GrLivArea alone
- Gives the variables more importance

Ridge			Lasso	
GrLivArea	0.29		GrLivArea	0.42
2ndFlrSF	0.26		OverallQual	0.18
KitchenQual_Gd	-0.21		KitchenQual_TA	-0.17
KitchenQual_TA	-0.20		KitchenQual_Gd	-0.16
LotArea	0.18		YearBuilt	0.13
YearBuilt	0.16		BsmtFinSF1	0.12
OverallQual	0.16		Neighborhood_NoRidge	0.09
ExterQual_Gd	-0.15		OverallCond	0.09
Exterior2nd_VinylSd	0.15		TotalBsmtSF	0.09
ExterQual_TA	-0.15		LotArea	0.08

Ridge Vs. Lasso

- ❑ Can optimize alpha for both regularizer
- ❑ Optimal mean test R^2 is better for LASSO
- ❑ Offers better feature selection

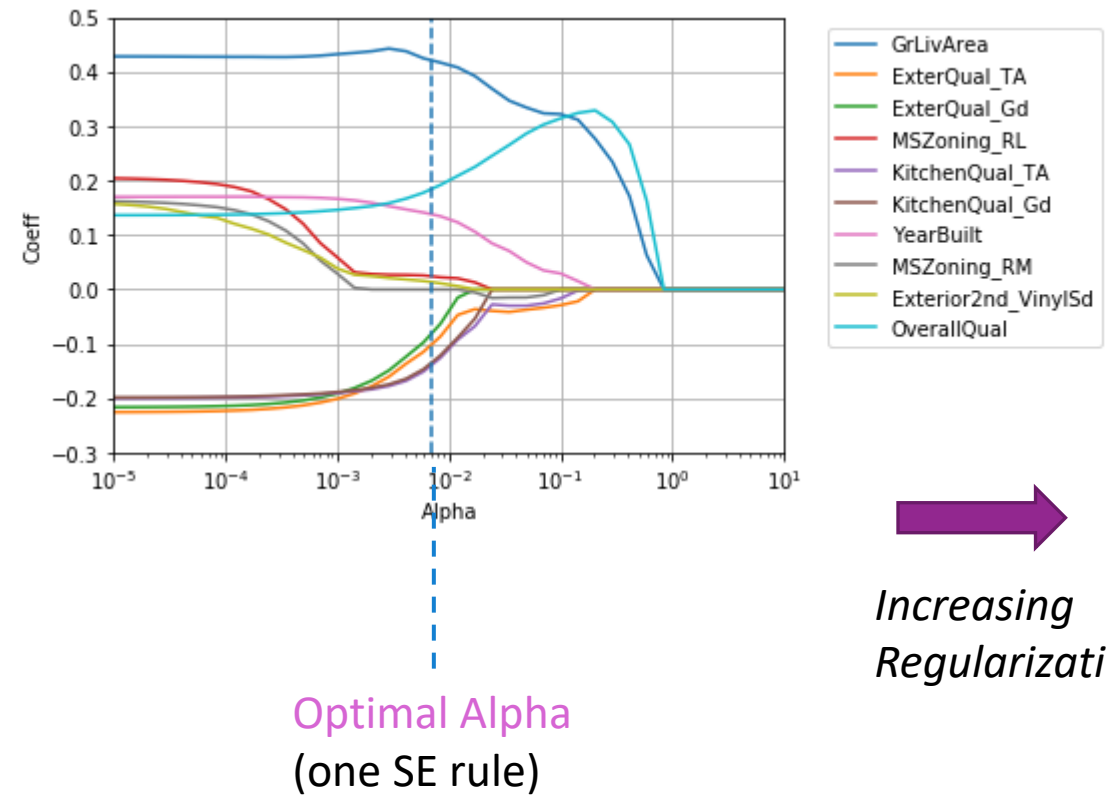


Optimal R^2 Lasso: 0.907283

Optimal R^2 Ridge: 0.897498

Lasso Path


- Plot of coefficients vs. alpha
- For large alpha:
 - All coefficients are zero
- As alpha is decreased:
 - One coefficient is activated at a time
 - Indicates an ordering of importance



Finding the Final Regressor

- ❑ Select features from cross-validation
- ❑ Re-run ordinary (un-regularized) regression on reduced features
- ❑ Use K –fold validation
- ❑ K -folds yield K weights and biases
- ❑ Take mean of the weights and biases for the final parameter estimate
- ❑ Take mean of the test MSE for the estimate of the test MSE

Outline

- ❑ Motivating Example: Feature selection in predicting housing prices
- ❑ Model selection and regularization
- ❑ Housing prices prediction with LASSO
- ❑ Probabilistic interpretation
- ❑ Other model selection examples
- ❑ Other model selection methods
- ❑ In-Class Exercise: Audio Pitch Detection

Maximum Likelihood Estimate

- Suppose that true data generated from probabilistic model with Gaussian noise:

$$\mathbf{y} = A\boldsymbol{\beta} + \mathbf{w}, \quad w_i \sim N(0, \sigma^2)$$

- Maximum likelihood estimator:

$$\hat{\boldsymbol{\beta}} = \arg \max_{\boldsymbol{\beta}} p(\mathbf{y}|A, \boldsymbol{\beta}) = \arg \min_{\boldsymbol{\beta}} [-\ln p(\mathbf{y}|A, \boldsymbol{\beta})]$$

- Gaussian density for noise in \mathbf{y} : $\ln p(\mathbf{y}|A, \boldsymbol{\beta}) = -\frac{1}{2\sigma^2} \|\mathbf{y} - A\boldsymbol{\beta}\|^2$

- Hence

$$\hat{\boldsymbol{\beta}} = \arg \max_{\boldsymbol{\beta}} p(\mathbf{y}|A, \boldsymbol{\beta}) = \arg \min_{\boldsymbol{\beta}} [\|\mathbf{y} - A\boldsymbol{\beta}\|^2] = \text{Least Squares Solution}$$

Bayes Estimation (MAP Estimate)

□ Maximum a posterior (MAP) estimator of β :

$$\hat{\beta} = \arg \max_{\beta} p(\beta | y, A)$$

- $\hat{\beta}$ = Most likely parameter value given evidence y, A

□ Bayes Rule: $p(\beta | y, A) = p(y|A, \beta)p(\beta)/p(y|A)$

□ Hence: $\hat{\beta} = \arg \max_{\beta} p(y|A, \beta)p(\beta)$ (because y and A are fixed)

- Likelihood: $p(y|A, \beta)$ How well β matches data
- Prior: $p(\beta)$: How well β agrees with prior knowledge about its distribution (constraints)

□ More in probability class...

Bayes Estimation with Logarithms

- Often easier to use logarithms:

$$\begin{aligned}\hat{\boldsymbol{\beta}} &= \arg \max_{\boldsymbol{\beta}} p(\mathbf{y} | A, \boldsymbol{\beta}) p(\boldsymbol{\beta}) = \arg \min_{\boldsymbol{\beta}} [-\ln p(\mathbf{y} | A, \boldsymbol{\beta}) p(\boldsymbol{\beta})] \\ &= \arg \min_{\boldsymbol{\beta}} [-\ln p(\mathbf{y} | A, \boldsymbol{\beta}) - \ln p(\boldsymbol{\beta})]\end{aligned}$$

- Gaussian density for noise in \mathbf{y} : $\ln p(\mathbf{y} | A, \boldsymbol{\beta}) = -\frac{1}{2\sigma^2} \|\mathbf{y} - A\boldsymbol{\beta}\|^2$

- Hence

$$\hat{\boldsymbol{\beta}} = \arg \min_{\boldsymbol{\beta}} \left[\frac{1}{2\sigma^2} \|\mathbf{y} - A\boldsymbol{\beta}\|^2 - \ln p(\boldsymbol{\beta}) \right] = \arg \min_{\boldsymbol{\beta}} [\|\mathbf{y} - A\boldsymbol{\beta}\|^2 + \phi(\boldsymbol{\beta})]$$

- **Conclusion:** MAP estimate = regularized LS with $\phi(\boldsymbol{\beta}) = -2\sigma^2 \ln p(\boldsymbol{\beta})$
 - Penalize $\boldsymbol{\beta}$ proportional to $-\ln p(\boldsymbol{\beta})$: Less likely $\boldsymbol{\beta}$ penalized more

Ridge and Lasso as Bayesian Estimators

- Bayesian Estimator:

$$\hat{\boldsymbol{\beta}} = \arg \min_{\boldsymbol{\beta}} \left[\frac{1}{2\sigma^2} \|\mathbf{y} - A \boldsymbol{\beta}\|^2 - \ln p(\boldsymbol{\beta}) \right]$$

- Assuming β_j are i.i.d. Gaussian with zero mean:

$$p(\beta_j) = \frac{1}{\sqrt{2\pi}\sigma} \exp(-\beta_j^2 / 2\sigma^2), \quad -\log p(\beta_j) = \beta_j^2 / 2\sigma^2 + \text{constants}$$


$$\hat{\boldsymbol{\beta}} = \arg \min_{\boldsymbol{\beta}} \left[\|\mathbf{y} - A \boldsymbol{\beta}\|^2 + \frac{\sigma^2}{\gamma^2} \|\boldsymbol{\beta}\|^2 \right] = \text{Ridge Regression!}$$

- Assuming β_j are i.i.d. Laplacian with zero mean:

$$p(\beta_j) = \frac{1}{2\sigma} \exp(-|\beta_j|/\sigma), \quad -\log p(\beta_j) = |\beta_j|/\sigma + \text{constant}$$

$$\hat{\boldsymbol{\beta}} = \arg \min_{\boldsymbol{\beta}} \left[\|\mathbf{y} - A \boldsymbol{\beta}\|^2 + \frac{2\sigma^2}{\gamma} \|\boldsymbol{\beta}\|_1 \right] = \text{Lasso Regression!}$$

Outline

- ❑ Motivating Example: Feature selection in predicting housing prices
- ❑ Model selection and regularization
- ❑ Housing prices prediction with LASSO
- ❑ Probabilistic interpretation
- ❑ Other model selection examples
- ❑ Other model selection methods
- ❑ In-Class Exercise: Audio Pitch Detection

Example 1: Medical Modeling

- ❑ Ex: Prostate specific antigen (PSA) test
 - Many years ago, PSA level was being consider for cancer screening
 - Question: Is a PSA test good for cancer?
 - Obtain features of prostate and correlate with PSA level
 - Determine if cancer volume is a relevant feature
 - See demo 1 on github site
 - Also in text

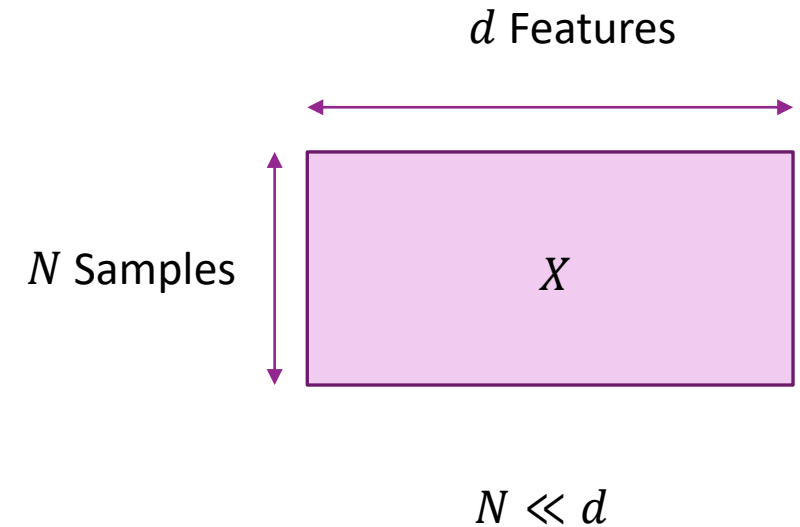


The data frame has the following components:

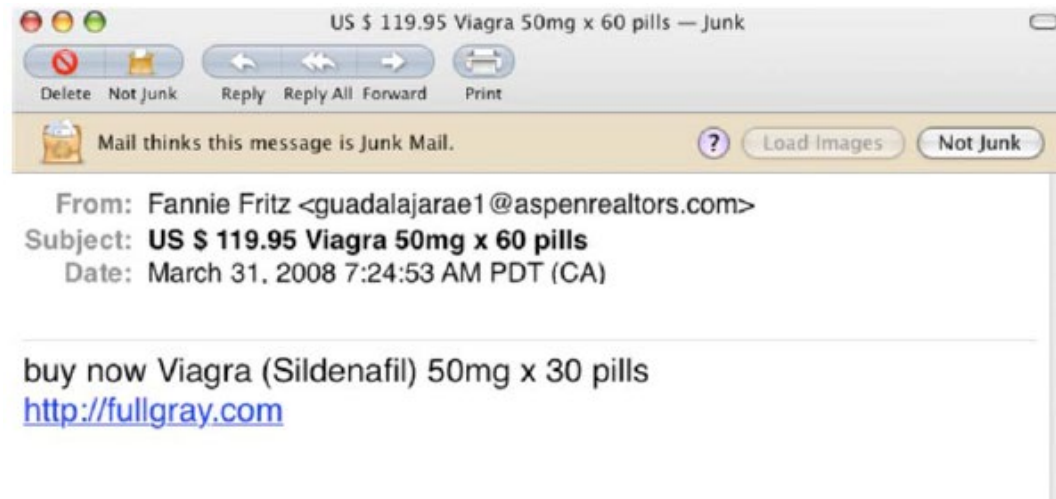
```
lcavol      log(cancer volume)
lweight     log(prostate weight)
age         age
lbph        log(benign prostatic hyperplasia amount)
svi         seminal vesicle invasion
lcp         log(capsular penetration)
gleason     Gleason score
pgg45       percentage Gleason scores 4 or 5
lpsa        log(prostate specific antigen)
```

Ex 2: Model Selection with Limited Data

- ❑ Model selection is particularly valuable when data is limited
- ❑ Ex: Consider linear model: $\hat{y} = b + w_1x_1 + \dots + w_dx_d$
 - Model has $d + 1$ parameters
- ❑ From previous lecture, we need $N > d + 1$ data points (x_i, y_i)
- ❑ In many cases we have $N \ll d$
 - Examples below
 - Many few data points than features
 - Classic linear fit will not work
- ❑ But, suppose we can restrict to $K \ll N$ non-zero parameters
 - Then, we can find a good fit on those parameters
- ❑ Challenge: How do we find a small number K of relevant features



Example 3: Spam Detection



❑ Classification problem:

- Is email junk or not junk?

❑ Typical bag-of-words model:

- Enumerate all words, $i = 1, \dots, d$
- Represent email via word count $x_i = \text{num instances of word } i$

❑ Model selection:

- d = vocabulary size is typically very large
- But, only a few words are likely relevant
- Want to find $K \ll d$ relevant words

Example 4: EEG

- ❑ EEG: Electroencephalography
- ❑ Measure brain activity from electrodes on scalp
- ❑ Source localization problem:
 - Find brain region responsible for evoked response
- ❑ Problem:
 - Many possible brain regions
Typically use $d > 10,000$ voxels
 - But, limited number of measurements:
100s of electrodes
 - Cannot fit a model from all brain regions
- ❑ Model selection:
 - We know that responses are likely from a small brain region
 - Find a small number of voxels that explain response
 - See lab!

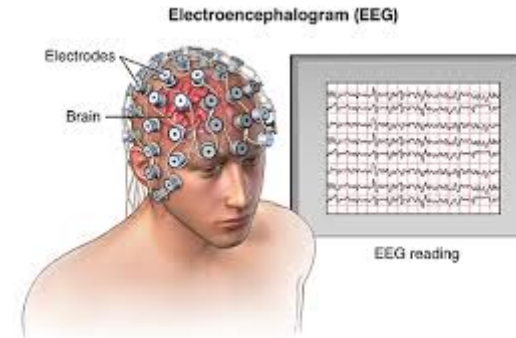

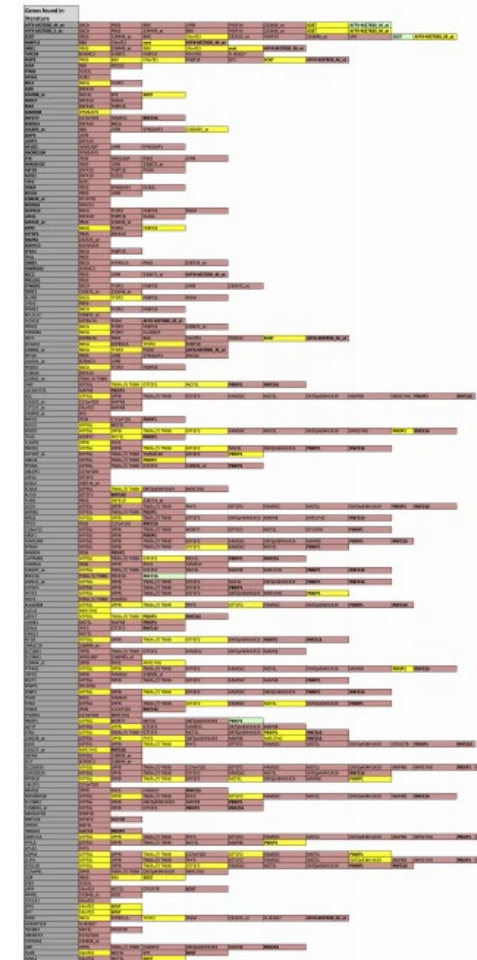


Image: mayoclinic.org




Example 5: DNA MicroArray Data

- ❑ Basic genetic problem
 - Which genes determine some characteristic (i.e. phenotype)?
- ❑ DNA microarrays:
 - Measure “expression” levels of large numbers of genes
 - Expression levels = amount of protein produced by gene
- ❑ Data modeling:
 - Fit phenotype to expression levels
 - Usually have large numbers of genes ($d \sim 1000$)
 - But, small number of data points ($n \sim 100$)
 - We know only a small number of genes are responsible
 - So, we can use model selection



Outline

- ❑ Motivating Example: Feature selection in predicting housing prices
- ❑ Model selection and regularization
- ❑ Housing prices prediction with LASSO
- ❑ Probabilistic interpretation
- ❑ Other model selection examples
- ❑ Other model selection methods
- ❑ In-Class Exercise: Audio Pitch Detection

Other feature selection methods

❑ Filtering method:

- Rank the features based on their correlation or mutual information with the target and possibly the redundancy among the features
- Simple but not very good

❑ Wrapper method:

- For each candidate feature subset, apply a chosen classifier/regressor, evaluate the cross validation accuracy. Go through all possible feature subsets, or test the subsets in some greedy way
- Computationally expensive

❑ Embedded method:

- Some regression/classification method naturally lead to feature ranking and selection

❑ What is available in Python:

- http://scikit-learn.org/stable/modules/feature_selection.html

Filtering method

- ❑ Rank the features based on their correlation with the target
 - Can use other metrics: Correlation, F-test, mutual information, ...
- ❑ Also should consider the redundancy (correlation) among chosen features
 - **Minimal Redundancy Maximum Relevance (mRMR)**
 - Peng, H.C., Long, F., and Ding, C., "Feature selection based on mutual information: criteria of max-dependency, max-relevance, and min-redundancy," IEEE Transactions on Pattern Analysis and Machine Intelligence, Vol. 27, No. 8, pp. 1226–1238, 2005.
 - <http://home.penglab.com/proj/mRMR/>
 - <https://www.mathworks.com/matlabcentral/fileexchange/14916-minimum-redundancy-maximum-relevance-feature-selection>

Ranking metrics

❑ Correlation coefficient between a feature and the target

❑ F-test: test the significance of using one feature vs. not using any (use the mean of y only. Essentially measure the difference in the MSE when using only the mean value of y vs. using a single feature.

$$f_{test} = \frac{r^2}{1-r^2}(n_{sample}-2)$$

❑ Mutual information between a feature and the target

$$I(X, Y) = \iint p(x, y) \log \frac{p(x, y)}{p(x)p(y)} dx dy$$

Embedded Method

- ❑ Results from some regression/classification methods allow feature selection
 - Linear regression: based on coefficient magnitude
 - Neural net: based on weight magnitude
 - Decision tree: based on tree level
 - Can add regularization terms on the coefficients/weights to encourage sparsity
 - LASSO regression
- ❑ Recursive feature elimination
 - Starting with all features, remove one feature that has the lowest importance (e.g. smallest coefficient magnitude)
 - Recursive feature elimination in sklearn
 - http://scikit-learn.org/stable/auto_examples/feature_selection/plot_rfe_digits.html#sphx-glr-auto-examples-feature-selection-plot-rfe-digits-py
 - http://scikit-learn.org/stable/auto_examples/feature_selection/plot_rfe_with_cross_validation.html#sphx-glr-auto-examples-feature-selection-plot-rfe-with-cross-validation-py

Wrapper method

- ❑ For each candidate feature subset, apply a chosen classifier/regressor, evaluate the cross validation accuracy. Go through all possible feature subsets, or test the subsets in some greedy way
 - Exhaustive search
 - Genetic algorithm
 - Forward stepwise
 - Backward stepwise

Exhaustive search for feature selection

- ❑ Suppose you want to consider feature subset of size up to p
- ❑ For all possible feature subsets of size 1 to p :
 - use cross validation to find mean RSS mean and standard deviation
- ❑ Choose the subset with the minimal RSS mean,
 - Or use the one standard error rule.
- ❑ When the number of features is large, may not be computationally feasible
- ❑ Fast search algorithms:
 - Genetic algorithm

Greedy feature selection

❑ Forward-Stepwise Selection

- Select one feature from all features that provides the lowest RSS with cross validation
- Select one new feature from all remaining features, so that previously chosen features plus the new feature provides the lowest RSS
- Repeat until the maximum feature number is reached, or when the RSS starts to increase

❑ Backward-Stepwise

- First use all features and find the RSS (using cross validation)
- Remove one feature and find the new RSS. Go through all possible features to remove.
- Find the one that leads to the least RSS increase. Remove this feature.
- Repeat the above, remove one from the remaining features, to find the next most important feature.

❑ Except exhaustive search, can all lead to suboptimal solution

Comparison of feature selection methods

Figure from [Hastie2008]: Hastie, Tibshirani, Friedman, The elements of statistical learning.

For more on this subject, see Sec. 3.3

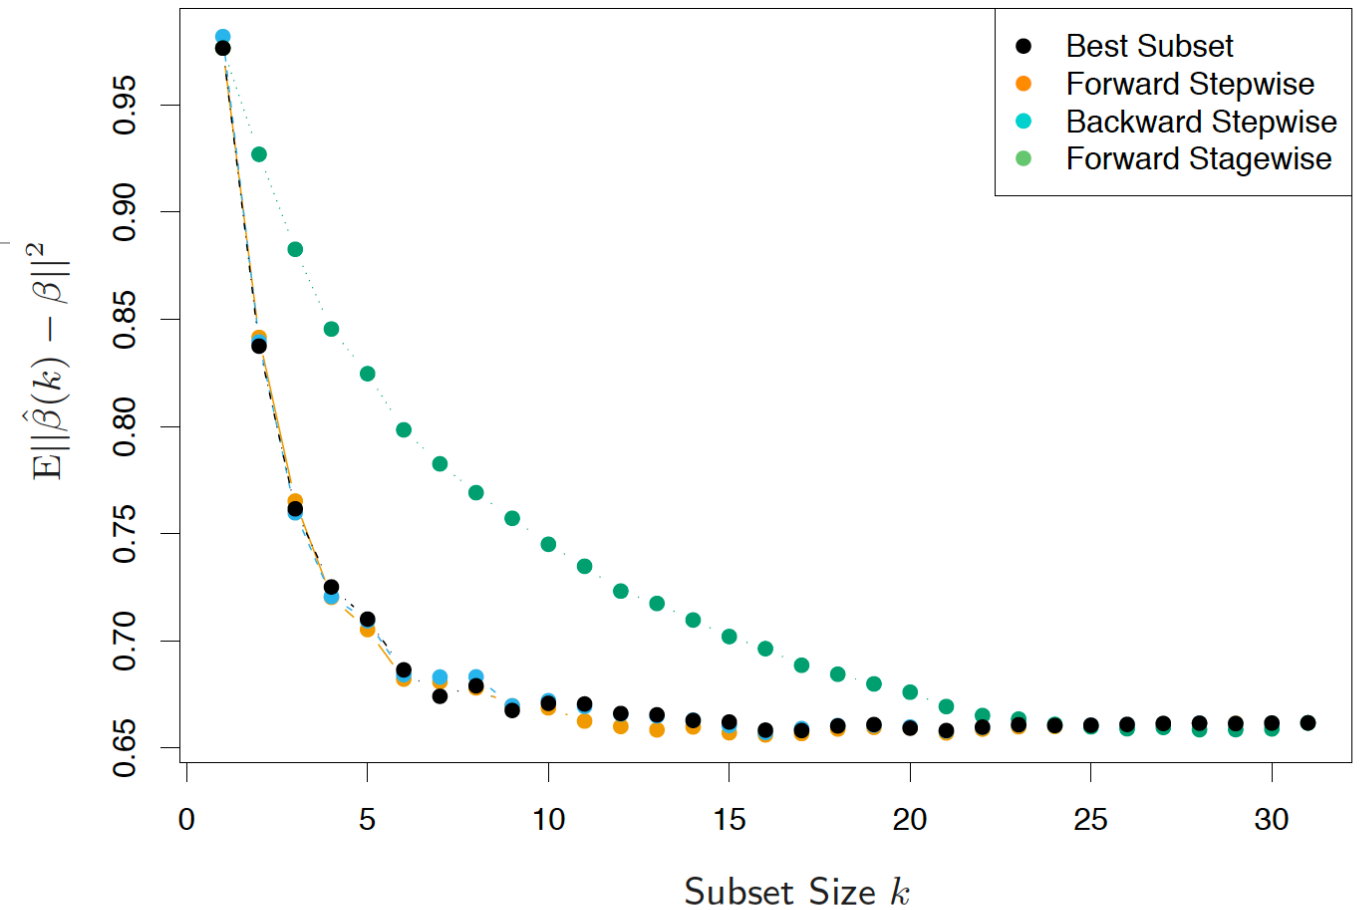


FIGURE 3.6. Comparison of four subset-selection techniques on a simulated linear regression problem $Y = X^T \beta + \varepsilon$. There are $N = 300$ observations on $p = 31$ standard Gaussian variables, with pairwise correlations all equal to 0.85. For 10 of the variables, the coefficients are drawn at random from a $N(0, 0.4)$ distribution; the rest are zero. The noise $\varepsilon \sim N(0, 6.25)$, resulting in a signal-to-noise ratio of 0.64. Results are averaged over 50 simulations. Shown is the mean-squared error of the estimated coefficient $\hat{\beta}(k)$ at each step from the true β .

More about cross validation

❑ Why do we use cross validation?

- To estimate the test error when there are insufficient training data so that we can partition the total data to a **large** training set and a **large** test set.
- Whether a dataset is large depends on the number of parameters of the model to be trained.
- Ideally the number of samples should be $>100x$ of the number of parameters, but at least $10x$.

❑ When you have sufficient training data, you can just use a certain percentage (e.g. 50%) for training and remaining for testing. The error on the testing set would be a reliable estimate of the test error.


❑ Two ways of using cross validation

- When the “best” model class, model order, and feature set are known:
 - Use CV to estimate the test error
- Use CV to determine the appropriate model class, model order and feature subset
 - For each candidate model class, model order, and feature subset, evaluate CV error
 - Determine which candidate yields the least CV error.

More about cross validation

- ❑ How to use the multiple estimated models from multiple trials?
 - Apply each on a test sample and take the average (for regression) or majority (for classification) of results
 - For linear regression, equivalent to average the model coefficients
- ❑ When your data is limited, you may want to go beyond K-folds
 - Ex: 5-fold means that you partition the data to 5 parts in some way, each part has 20% of data, and only do 5 fold training and testing
 - When your data is small, the average CV error is still very sensitive to how the data is partitioned to 5 parts. If you use random shuffling, you will get different result each time.
 - Instead, you could do L trials ($L \gg 5$) of random sampling, each time using 80% for training and 20% for testing
- ❑ How to handle limited data in machine learning is still a challenging topic!

Outline

- ❑ Motivating Example: Predicting prostate cancer from a PSA test
- ❑ Model Selection
- ❑ Model Selection from LASSO regularization
- ❑ Probabilistic interpretation
- ❑ Other Model Selection Methods
- ❑ In-Class Exercise: Audio Pitch Detection

In Class Exercise

☐ https://github.com/sdrangan/introml/blob/master/unit05_lasso/lasso_in_class.ipynb

LASSO Regression In-Class Exercise

In this exercise, we will see how to use LASSO for pitch detection in audio.

We load the following packages.

```
import numpy as np
import matplotlib.pyplot as plt
import pickle
```

Load the data

The data is taken from a sample of about 20~ms of audio from a viola. I have already pre-processed the data. You can load it with following command. The value t is the time (in seconds) and y is the sample of audio (this is a mono recording).

```
fn_src = 'https://raw.githubusercontent.com/sdrangan/introml/master/unit05_lasso/viola_sample.p'
fn_dst = 'viola_sample.p'

import os
from six.moves import urllib

if os.path.isfile(fn_dst):
```

What You Should Know to Do

- ☐ Formulate a linear estimation problem with a regularization
- ☐ Compute an L1-regularized estimate (LASSO) using sklearn tools
- ☐ Compute the optimal regularization level using cross validation
- ☐ Interpret results from a LASSO path
- ☐ Determine final regression function from cross validation
- ☐ Set regularizer based on a probabilistic prior
- ☐ Perform other feature selection methods