

Introduction to Machine Learning

Problems: LASSO and Model Selection

Prof. Sundeep Rangan

1. *Exhaustive search.* In this problem, we will look at how to exhaustively search over all possible subsets of features. You are given three python functions:

```
model = LinearRegression() # Create a linear regression model object
model.fit(X,y)              # Fits the model
yhat = model.predict(X)     # Predicts targets given features
```

Given training data $\mathbf{x}_{tr}, \mathbf{y}_{tr}$ and test data $\mathbf{x}_{ts}, \mathbf{y}_{ts}$, write a few lines of python code to:

- Find the best model using only one feature of the data (i.e. one column of \mathbf{x}_{tr} and \mathbf{x}_{ts}).
- Find the best model using only two features of the data (i.e. two columns of \mathbf{x}_{tr} and \mathbf{x}_{ts}).
- Suppose we wish to find the best k of p features via exhaustive searching over all possible subsets of features. How many times would you need to call the `fit` function? What if $k = 10$ and $p = 1000$?

2. *Selecting a regularizer.* Suppose we fit a regularized least squares objective,

$$J(\mathbf{w}) = \sum_{i=1}^N (y_i - \hat{y}_i)^2 + \lambda \phi(\mathbf{w}),$$

where \hat{y}_i is some prediction of y_i given the model parameters \mathbf{w} . For each case below, suggest a possible regularization function $\phi(\mathbf{w})$. There is no single correct answer.

- All parameters vectors \mathbf{w} should be considered.
- Negative values of w_j are unlikely (but still possible).
- For each j , w_j should not change that significantly from w_{j-1} .
- For most j , $w_j = w_{j-1}$. However, it can happen that w_j can be different from w_{j-1} for a few indices j .

Variable	Units	Mean	Std dev
Median income, x_1	\$	50000	15000
Median age, x_2	years	45	10
House sale price, y	\$1000	300	100

Table 1: Features for Problem 3

3. *Normalization.* A data analyst for a real estate firm wants to predict house prices based on two features in each zip code. The features are shown in Table 1. The agent decides to use a linear model,

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2, \quad (1)$$

- (a) What is the problem in using a LASSO regularizer of the form,

$$\phi(\beta) = \sum_{j=1}^2 |\beta_j|.$$

- (b) To uniformly regularize the features, she fits a model on the normalized features,

$$\hat{u} = \alpha_1 z_1 + \alpha_2 z_2, \quad z_j = \frac{x_j - \bar{x}_j}{s_j}, \quad u = \frac{\hat{y} - \bar{y}}{s_y},$$

where s_j and s_y are the standard deviations of the x_j and y . She obtains parameters $\alpha = [0.6, -0.3]$? What are the parameters β in the original model (1)?

4. *Normalization in python.* You are given python functions,

```
model = SomeModel()           # Creates a model
model.fit(Z,u)                 # Fits the model, expecting normalized features
yhat = model.predict(Z)       # Predicts targets given features
```

Given training data $\mathbf{x}_{tr}, \mathbf{y}_{tr}$ and test data $\mathbf{x}_{ts}, \mathbf{y}_{ts}$, write python code to:

- Normalize the training data to remove the mean and standard deviation from both \mathbf{x}_{tr} and \mathbf{y}_{tr} .
- Fit the model on the normalized data.
- Predict the values \mathbf{y}_{hat} on the test data.
- Measure the RSS on the test data.

5. *Discretization.* Suppose we wish to fit a model,

$$y \approx \hat{y} = \sum_{j=1}^K \beta_j e^{-\alpha_j x}, \quad (2)$$

for parameters α_j and β_j . Since the parameters α_j are not known, this model is nonlinear and cannot be fit with least squares. A common approach in such circumstances is to use an alternate linear model,

$$y \approx \hat{y} = \sum_{j=1}^p \tilde{\beta}_j e^{-\tilde{\alpha}_j x}, \quad (3)$$

where the values $\tilde{\alpha}_1, \dots, \tilde{\alpha}_p$ are a *fixed*, large set of possible values for α_j , and $\tilde{\beta}_j$ are the coefficients in the model. Since the values $\tilde{\alpha}_j$ are fixed, only the parameters $\tilde{\beta}_j$ need to be learned. Hence, the model (3) is linear. The model (3) is equivalent to (2) if only a small number K of the coefficients $\tilde{\beta}_j$ are non-zero. You are given three python functions:

```

model = Lasso(lam=lam)           # Creates a linear LASSO model
                                  # with a regularization lam
beta = model.fit(Z,y)            # Finds the model parameters using the
                                  # LASSO objective
                                  # ||y-Z*beta||^2 + lam*||beta||_1
yhat = model.predict(Z)          # Predicts targets given features Z:
                                  # yhat = Z*beta

```

Note this syntax is slightly different from the `sklearn` syntax. You are also given training data `xtr,ytr` and test data `xts,yts`. Write python code to:

- Create $p = 100$ values of $\tilde{\alpha}_j$ uniformly in some interval $\tilde{\alpha}_j \in [a, b]$ where a and b are given.
 - Fit the linear model (3) on the training data for some given `lam`.
 - Measure the test error.
 - Find coefficients α_j and β_j corresponding to the largest $k = 3$ values in $\tilde{\beta}_j$. You can use the function `np.argsort`.
6. *Minimizing an ℓ_1 objective.* In this problem, we will show how to minimize a simple scalar function with an ℓ_1 -term. Given y and $\lambda > 0$, suppose we wish to find the minimum,

$$\hat{w} = \arg \min_w J(w) = \frac{1}{2}(y - w)^2 + \lambda|w|.$$

Write \hat{w} in terms of y and λ . Since $|w|$ is not differentiable everywhere, you cannot simply set $J'(w) = 0$ and solve for w . Instead, you have to look at three cases:

- (i) First, suppose there is a minima at $w > 0$. In this region, $|w| = w$. Since the set $w > 0$ is open, at any minima $J'(w) = 0$. Solve for w and test if the solution indeed satisfies $w > 0$.
- (ii) Similarly, suppose $w < 0$. Solve for $J'(w) = 0$ and test if the solution satisfies the assumption that $w < 0$.
- (iii) If neither of the above cases have a minima, then the minima must be at $w = 0$.